Foundations of Machine Learning

Brett Bernstein

August 22, 2018

EM: Concept Check

EM/Mixture Model Objectives

- Write down the probability model corresponding to the GMM problem (multinomial distribution on mixture component z, multivariate Gaussian conditionals $x|z$).
- Write down the joint density $p(x, z)$ for the GMM model.
- Give an expression for the marginal log-likelihood for the observed data x for the GMM model, and explain why it doesn't simplify as nicely as the log-likelihood of a multivariate Gaussian model.
- Give pseudocode for the EM Algorithm for GMM (as in slide 29).
- Give an expression for the probability model for a generic latent variable model, where x is observed, z is latent (i.e. unobserved), and the parameters are represented by θ.
- Give EM algorithm pseudocode (as in slide 27).

EM Question

Poisson Mixture Model Setup: Consider the poisson mixture model, where each data instance is generated as follows:

1. Draw an [unobserved] cluster assignment z from a multinomial distribution $\pi = (\pi_1, \cdots, \pi_k)$ on k clusters.
2. Draw a count from a Poisson distribution with PMF:

$$p(x \mid \lambda_z) = \frac{\lambda_z^x e^{-\lambda_z}}{x!},$$

where $\lambda = (\lambda_1, \ldots, \lambda_k) \in (0, \infty)^k$.

To keep things concise, we'll write $\theta = (\pi, \lambda)$ to represent all of the unknown parameters.

Problems:

1. To start, let x, z be the count and cluster assignment for a single instance. Give an expression for $p(x, z \mid \theta)$ in terms of $p(z \mid \theta)$ and $p(x|z, \theta)$.

2. Give an expression for $p(x \mid \theta)$, the marginal distribution for a single observed x, in terms of π, λ, and x.

3. Give an expression for the conditional distribution $p(z \mid x, \theta)$ in terms of $p(x, z \mid \theta)$ and $p(x \mid \theta)$. (Basic probability review)

4. Now assume we have some training set of size n. We observe $x = (x_1, \cdots, x_n)$, but don't observe $z = (z_1, \cdots, z_n)$. We'll work through the EM algorithm for this problem. First, let's tackle the "E step", in which we evaluate the responsibilities $\gamma_i^j = p(z_i = j|x_i)$ for each $j \in \{1, \cdots, k\}$. Give an expression for this responsibility for cluster j and instance i.

5. Before we move on to the "M step", let's apply this "E step" result to a toy problem. Imagine $k = 3$, and we have $\lambda_1 = 1$, $\lambda_2 = 2$, and $\lambda_3 = 3$. Find $p(z = 2|x = 1)$ in terms of π_i for i in $\{1, 2, 3\}$. Hint: Note $p(x)$ is constant for all k, so its straightforward to give proportional expressions for each of $p(z = k|x = 1)$ then normalize.

6. Now we will tackle the "M step" of the EM algorithm, during which we will update π_z and λ_z. To start, find our objective (the expectation of the complete log-likelihood).

7. Finally give the expression for λ_z^{new} (that maximizes the objective you found above).